

# End-to-End ML Project

## – California Housing Price Prediction

Gauranga Kumar Baishya

August 27, 2025

# Outline

- 1 The Big Picture
- 2 Getting the Data
- 3 Discover & Visualize Data
- 4 Data Preparation
- 5 Select & Train a Model
- 6 Fine-Tune Your Model
- 7 Launch, Monitor, & Maintain

# Project Goal & The Data Pipeline

## Primary Objective

To build a model that predicts the median housing price in California districts using 1990 census data.

## Business Context

The model's output will be a crucial **signal** for a downstream investment analysis system. The accuracy of this prediction directly impacts business revenue.

## Data Pipeline

This project involves creating a full data pipeline: a sequence of data processing components, from raw data to the final actionable insight.

# Framing the Problem

## 1. Learning Type

**Supervised Learning:** The dataset includes the target labels (median housing prices).

## 2. Task Type

**Regression Task:** We are predicting a continuous numerical value. Specifically, it is a **multivariate regression** problem as we use multiple features.

## 3. Learning Method

**Batch Learning:** The dataset is small enough to be handled in memory, and there is no need for real-time adaptation.

# Selecting a Performance Measure

## Root Mean Square Error (RMSE)

The most common metric for regression tasks. It measures the standard deviation of the prediction errors and gives a higher weight to large errors.

$$\text{RMSE}(\mathbf{X}, h) = \sqrt{\frac{1}{m} \sum_{i=1}^m (h(\mathbf{x}^{(i)}) - y^{(i)})^2}$$

RMSE is sensitive to outliers because of the squaring operation ( $l_2$  norm).

## Mean Absolute Error (MAE)

An alternative metric; MAE is less sensitive to outliers..

$$\text{MAE}(\mathbf{X}, h) = \frac{1}{m} \sum_{i=1}^m |h(\mathbf{x}^{(i)}) - y^{(i)}|$$

# How does the data look like ?

Summary Statistics									
	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value
count	20640.000000	20640.000000	20640.000000	20640.000000	20433.000000	20640.000000	20640.000000	20640.000000	20640.000000
mean	-119.569704	35.631861	28.639486	2635.763081	537.870553	1425.476744	499.539680	3.870671	206855.816909
std	2.003532	2.135952	12.585558	2181.615252	421.385070	1132.462122	382.329753	1.899822	115395.615874
min	-124.350000	32.540000	1.000000	2.000000	1.000000	3.000000	1.000000	0.499900	14999.000000
25%	-121.800000	33.930000	18.000000	1447.750000	296.000000	787.000000	280.000000	2.563400	119600.000000
50%	-118.490000	34.260000	29.000000	2127.000000	435.000000	1166.000000	409.000000	3.534800	179700.000000
75%	-118.010000	37.710000	37.000000	3148.000000	647.000000	1725.000000	605.000000	4.743250	264725.000000
max	-114.310000	41.950000	52.000000	39320.000000	6445.000000	35682.000000	6082.000000	15.000100	500001.000000

# How does the data look like ?




Figure 2-8. A histogram for each numerical attribute

# Initial Data Exploration with Pandas

## Key Commands

- : Reveals dataset size (e.g., 20,640 entries), data types, and missing values.
- : Shows that `ocean_proximity` is a categorical feature with 5 distinct categories.
- : Provides a statistical summary of all numerical attributes.
- : Plots histograms to visualize the distribution of each numerical feature.



# Key Findings from Histograms

- 1 **Capped Data:** `housing_median_age` and `median_house_value` are capped. The cap on the target variable at \$500,000 is a significant issue.
- 2 **Scaled Income:** `median_income` is not in USD but is scaled and capped at 15.
- 3 **Varying Scales:** Attributes have vastly different scales, indicating a need for **feature scaling**.
- 4 **Tail-Heavy Distributions:** Many attributes are not normally distributed. They may need transformation (e.g., log transform) to help algorithms detect patterns.

# Creating a Test Set

## Why is this CRITICAL?

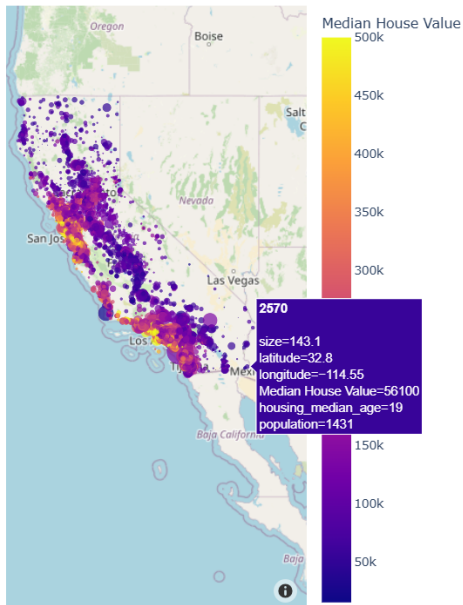
You must create and set aside a test set **before** any deep data exploration. This prevents **data snooping bias**, where you might unconsciously select a model based on patterns you see in the test set, leading to an overly optimistic performance evaluation.

## Best Practice: Stratified Sampling

Instead of pure random sampling, **stratified sampling** is used to ensure the test set is representative.

- Since median income is a strong predictor, we stratify based on income categories.
- This creates a test set where the distribution of income categories matches the distribution in the full dataset, avoiding sampling bias.

# Geographical Visualization & Correlation



# Geographical Visualization & Correlation




Figure 2-15. Scatter matrix

# Geographical Visualization & Correlation

## Geographical Plot

A scatterplot of latitude vs. longitude, with point density ('alpha'), population size ('s'), and price ('c'), reveals that prices are highest in dense, coastal areas like the Bay Area and LA.

## Correlation Analysis

Using `.corr()`, we find `median_income` has the strongest positive correlation with house value. The `scatter_matrix` confirms this and reveals data quirks, like horizontal lines at price caps.

## Attribute Combinations

Creating new features can be very powerful. For example:

- `rooms_per_household`
- `bedrooms_per_room`

These new features proved to be more correlated with house value than the original attributes.

# Cleaning, Scaling, and Pipelines

## Handling Missing Values

The missing `total_bedrooms` values are handled using Scikit-Learn's `Imputer`, configured to fill them with the column's median.

## Handling Categorical Data

The text attribute `ocean_proximity` is converted to numerical data using **One-Hot Encoding**, which creates a binary column for each category. This prevents algorithms from assuming a false ordering.

## Feature Scaling

- **Normalization** (`MinMaxScaler`): Scales to a 0-1 range. Sensitive to outliers.
- **Standardization** (`StandardScaler`): Centers to mean 0 with unit variance. More robust to outliers.

# Training and Evaluating Models

## Model 1: Linear Regression

- A simple baseline model.
- **Result:** High RMSE on the training set,  $\approx 68.628$ .
- **Diagnosis:** The model is **underfitting**. It's too simple to capture the data's complexity.

## Model 2: Decision Tree Regressor

- A more powerful, non-linear model.
- **Result:** RMSE of 0 on the training set.
- **Diagnosis:** A classic case of severe **overfitting**. The model has memorized the training data and will not generalize.

# Better Evaluation: Cross-Validation

## K-Fold Cross-Validation

To get a more robust performance estimate without touching the test set, we use K-fold cross-validation. The training set is split into K folds; the model is trained K times on K-1 folds and evaluated on the remaining fold.

## Cross-Validation Results

- **Decision Tree:** The average CV score was worse than Linear Regression, confirming it was overfitting.
- **Random Forest Regressor:** An ensemble model that trains many Decision Trees. It performed much better, with a CV RMSE of  $\approx 52.564$ . This is our most promising model.



# Hyperparameter Tuning

## Grid Search

Scikit-Learn's `GridSearchCV` automates the process of finding the best hyperparameters.

- You define a "grid" of hyperparameters and values to test.
- It uses cross-validation to evaluate every combination.
- The search improved the Random Forest RMSE to  $\approx 49.694$ .

## Analyze Feature Importance

The tuned Random Forest model can report **feature importances**.

- **Most important:** `median_income`.
- **Second most important:** The categorical feature `INLAND`.
- This insight could be used to drop less important features.

# Final Evaluation on the Test Set

## The Moment of Truth

After all model selection and tuning is complete, the final model is evaluated one last time on the held-out test set.

## Final Performance

The final RMSE on the test set was  $\approx 47,766$ .

## Important Rule

You must **not** perform any further tuning based on the test set performance. This result represents the model's expected performance on new, unseen data.

# Deployment and Beyond

## Launch

Deploy the final model into your production environment, connecting it to live data sources.

## Monitor

- Continuously monitor the model's live performance.
- Models can degrade over time as data evolves—a concept known as "**model rot**".
- Set up alerts to be notified of performance drops.

## Maintain

- Automate the process of retraining your model on fresh data.
- Regular retraining ensures the model stays relevant and its performance remains high.

# Thank You!